

ADS 506 - Final Project EDA

Emanuel Lucban, Sean Torres, Christopher Richardson

```
library(astsa)
library(caret)
library(corrplot)
library(DataExplorer)
library(DMwR)
library(dplyr)
library(ggplot2)
library(forecast)
library(neuralnet)
library(tidyr)
library(tidyverse)

set.seed(1)
```

Loading Data

DF summary/information and dimensions

```
df <- read_csv('Data/water_quality_2011_2019_datasd.csv')
summary(df)
```

```
##      sample      station      depth_m      date_sample
## Min.   :1.011e+08 Length:338978 Min.    : 1.00 Min.    :2011-01-01
## 1st Qu.:1.301e+09 Class :character 1st Qu.: 2.00 1st Qu.:2013-01-23
## Median :1.505e+09 Mode  :character Median : 9.00 Median :2015-05-08
## Mean   :1.416e+09      Mean  :13.18 Mean  :2015-05-10
## 3rd Qu.:1.708e+09      3rd Qu.:18.00 3rd Qu.:2017-08-02
## Max.   :1.912e+09      Max.   :98.00 Max.   :2019-12-30
##              NA's   :29894
##      time      project      parameter      qualifier
## Length:338978 Length:338978 Length:338978 Length:338978
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##      value      units
## Min.   :    0.00 Length:338978
## 1st Qu.:    2.00 Class :character
## Median :    8.17 Mode  :character
## Mean   :   77.16
## 3rd Qu.:   25.07
## Max.   :180000.00
```

```
## NA's :670
```

```
dim(df)
```

```
## [1] 338978 10
```

Station Count

```
length <- length(unique(df$station))
line <- paste('There is a total of', length, 'stations!')
cat(line, '\n\nHere are the stations from the dataset.')
```

```
## There is a total of 105 stations!
```

```
##
```

```
## Here are the stations from the dataset.
```

```
sort(unique(df$station))
```

```
## [1] "A1" "A6" "A7" "C4" "C5" "C6" "C7" "C8" "D10" "D11"
## [11] "D12" "D4" "D5" "D7" "D8" "D8-A" "D8-B" "D9" "F01" "F02"
## [21] "F03" "F04" "F05" "F06" "F07" "F08" "F09" "F10" "F11" "F12"
## [31] "F13" "F14" "F15" "F16" "F17" "F18" "F19" "F20" "F21" "F22"
## [41] "F23" "F24" "F25" "F26" "F27" "F28" "F29" "F30" "F31" "F32"
## [51] "F33" "F34" "F35" "F36" "I1" "I10" "I11" "I12" "I13" "I14"
## [61] "I15" "I16" "I17" "I18" "I19" "I2" "I20" "I21" "I22" "I23"
## [71] "I24" "I25" "I26" "I27" "I28" "I29" "I3" "I30" "I31" "I32"
## [81] "I33" "I34" "I35" "I36" "I37" "I38" "I39" "I4" "I40" "I5"
## [91] "I6" "I7" "I8" "I9" "S0" "S10" "S11" "S12" "S2" "S3"
## [101] "S4" "S5" "S6" "S8" "S9"
```

Parameter Count

```
length <- length(unique(df$parameter))
line <- paste('There is a total of', length, 'parameters!')
cat(line, '\n\nHere are the parameters from the dataset.')
```

```
## There is a total of 12 parameters!
```

```
##
```

```
## Here are the parameters from the dataset.
```

```
sort(unique(df$parameter))
```

```
## [1] "CHLOROPHYLL" "DENSITY" "DO" "ENTERO" "FECAL"
## [6] "OG" "PH" "SALINITY" "SUSO" "TEMP"
## [11] "TOTAL" "XMS"
```

Kelp Stations and Parameters

```
# We are only interested in kelp stations (total count = 7)
kelp_stations <- c("I19", "I24", "I25", "I26", "I32", "I39", "I40")
```

```
# Variables/Parameters utilized
parameters <- c("CHLOROPHYLL",
               "DO",
               "ENTERO",
               "FECAL",
```

```
"PH",
"SALINITY",
"TEMP")
```

```
df <- df[df$parameter %in% parameters,]
df <- df[df$station %in% kelp_stations,]
```

In total we only want the 7 kelp stations within the data set as well as the 7 types of measurements.


```
# Remove individual Date & Time features + project since all buoys are related to SB00
columns2drop <- c('qualifier', 'project', 'sample', 'time')
df <- drop_columns(df, columns2drop)
head(df,3)
```

```
## # A tibble: 3 x 6
##   station depth_m date_sample parameter value units
##   <chr>      <dbl> <date>      <chr>      <dbl> <chr>
## 1 I25          2 2011-01-01 ENTERO         24 CFU/100 mL
## 2 I25          6 2011-01-01 ENTERO        110 CFU/100 mL
## 3 I25          9 2011-01-01 ENTERO        100 CFU/100 mL
```

Null Ratios by Station

```
for (s in kelp_stations){
  bouy <- df[df$station == s,]

  plot_missing(bouy, title=paste('station', s))
}
```



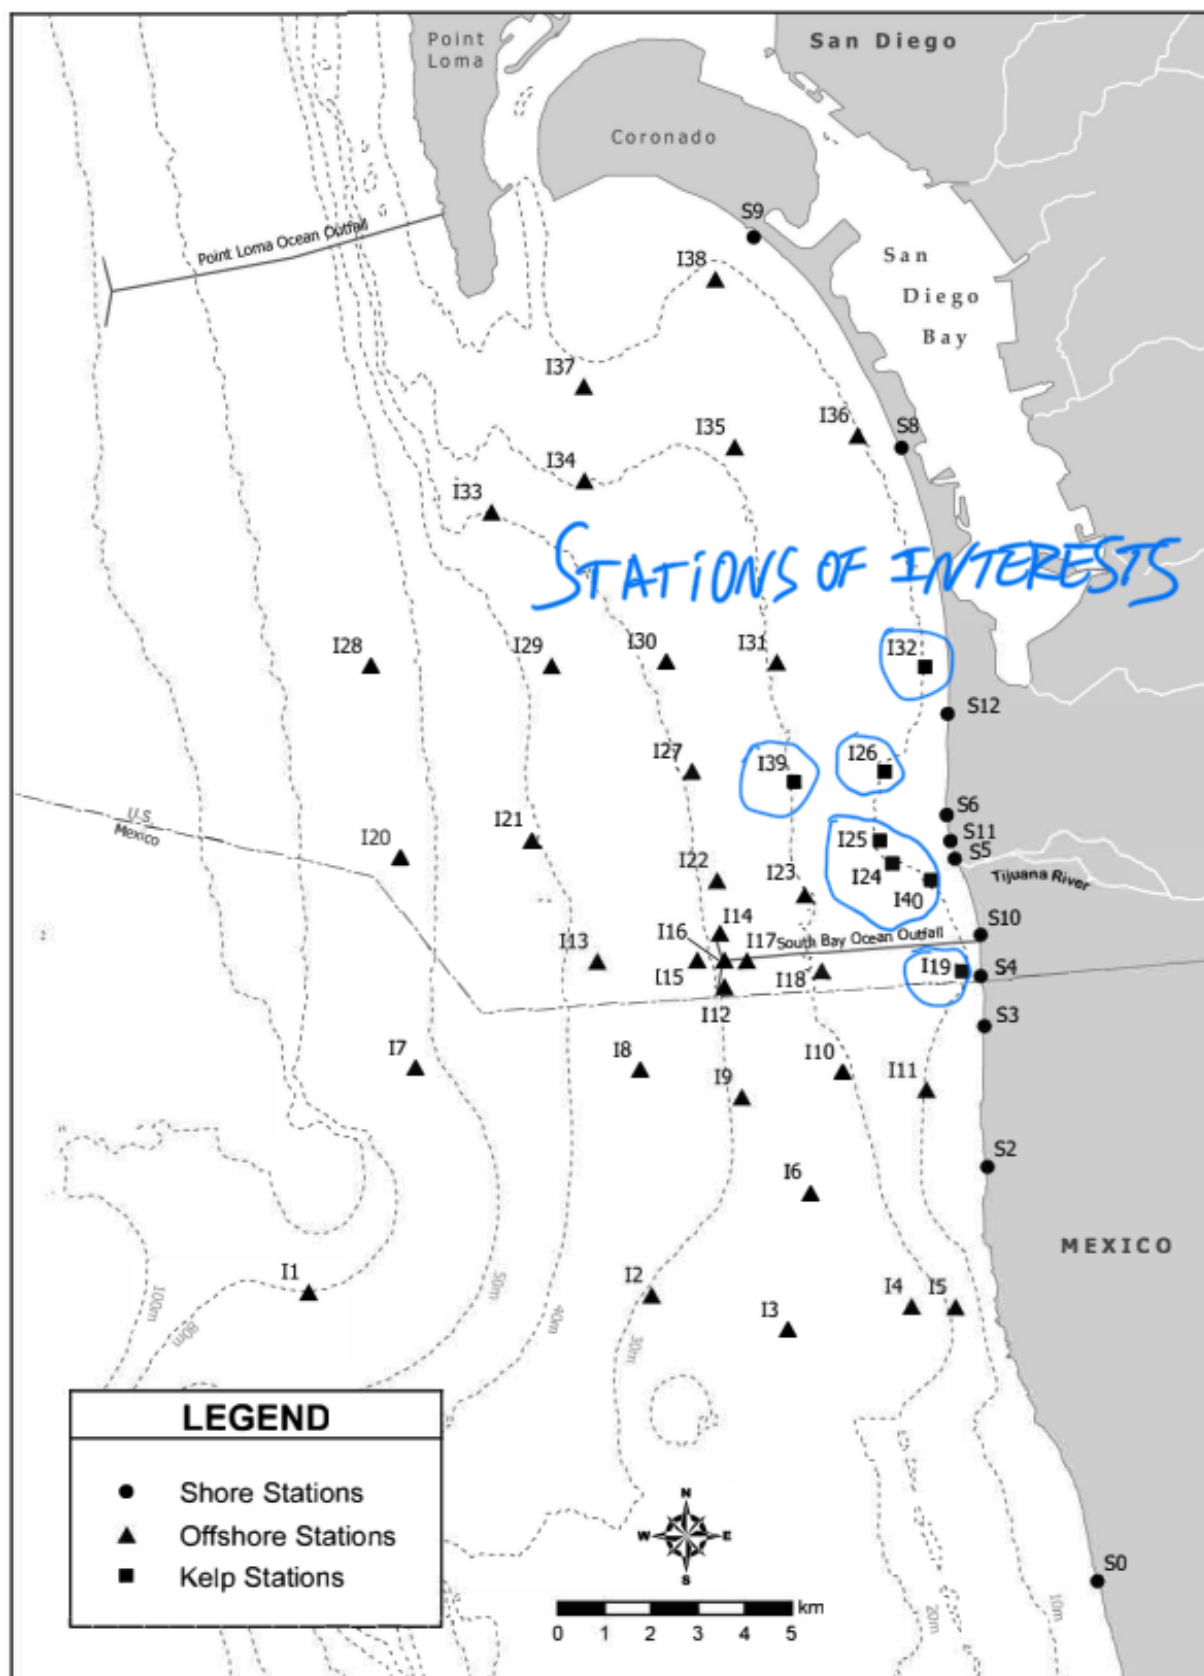


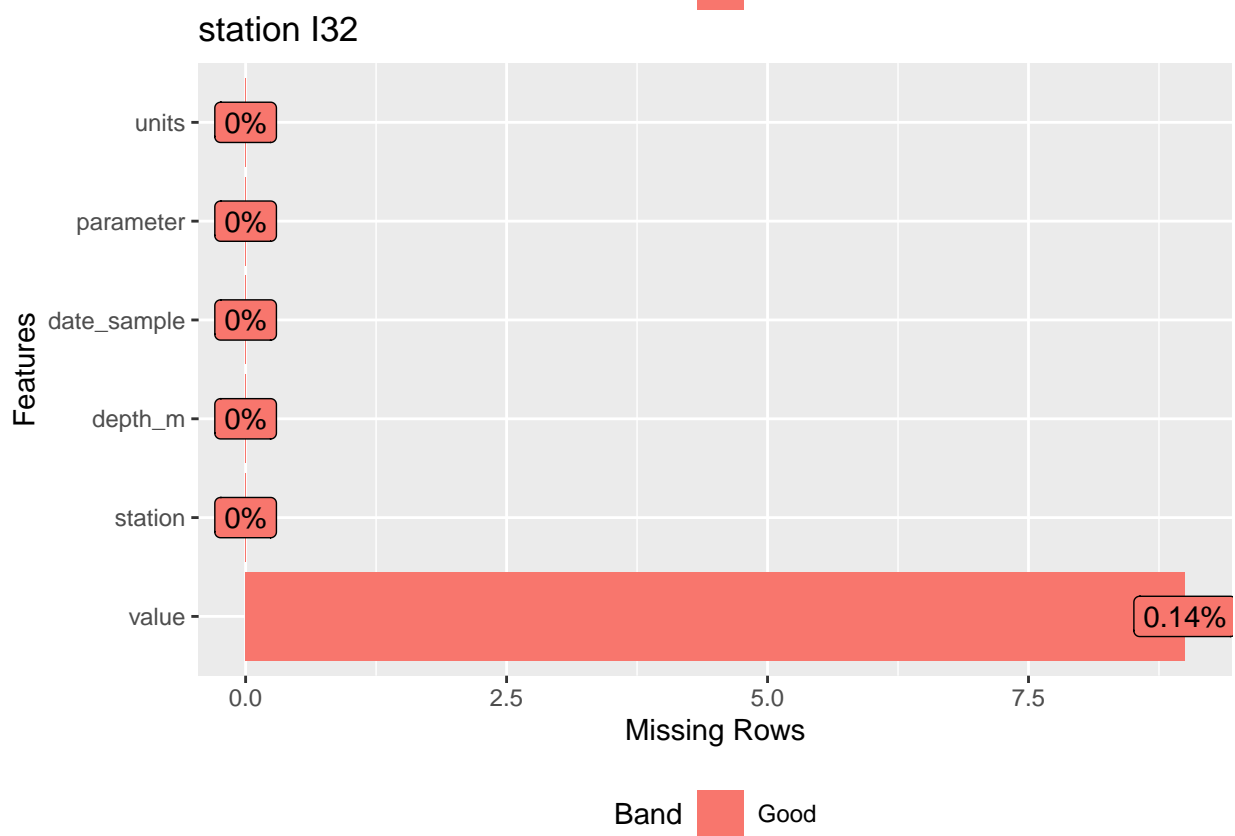


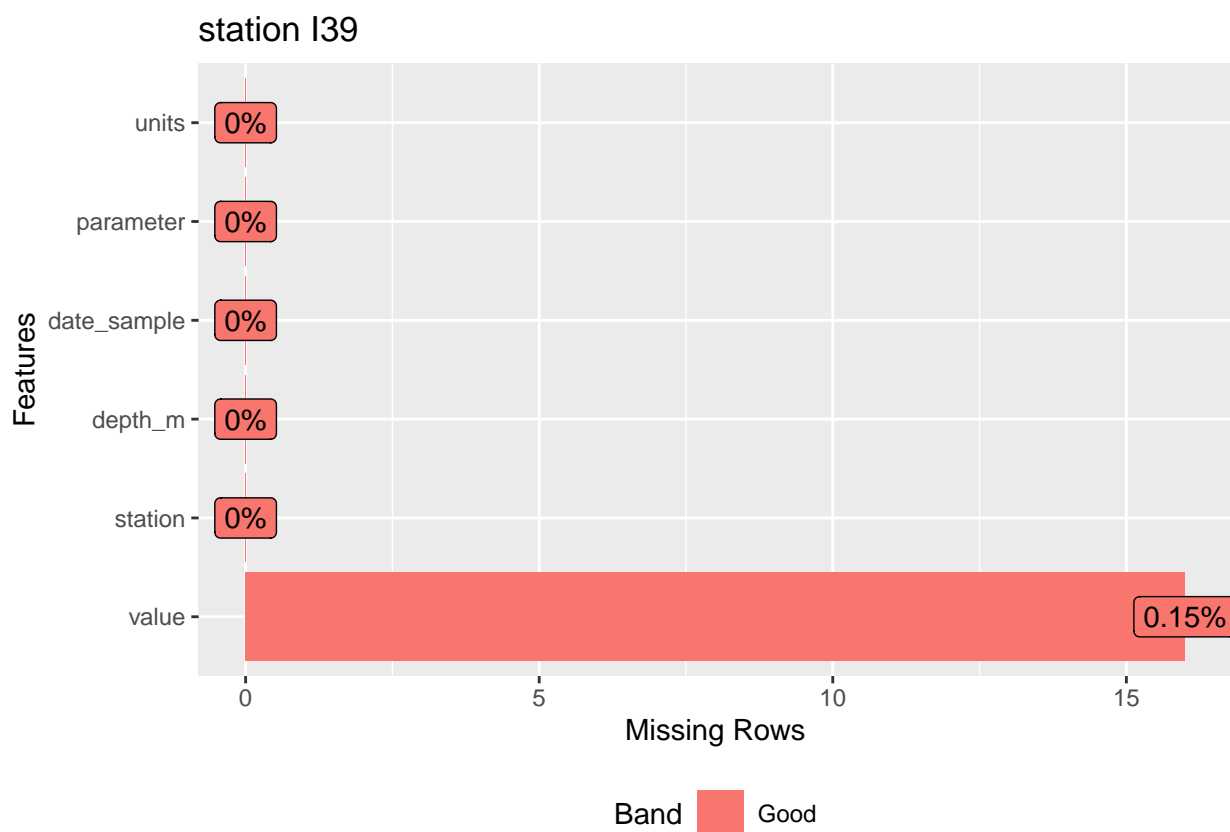


Figure 1: Map of interests








```



# Removal of null values
df <- df[!is.na(df$value),]


# validation of dropping of NA rows
for (s in kelp_stations){
  bouy <- df[df$station == s,]

  plot_missing(bouy, title=paste('station', s))
}

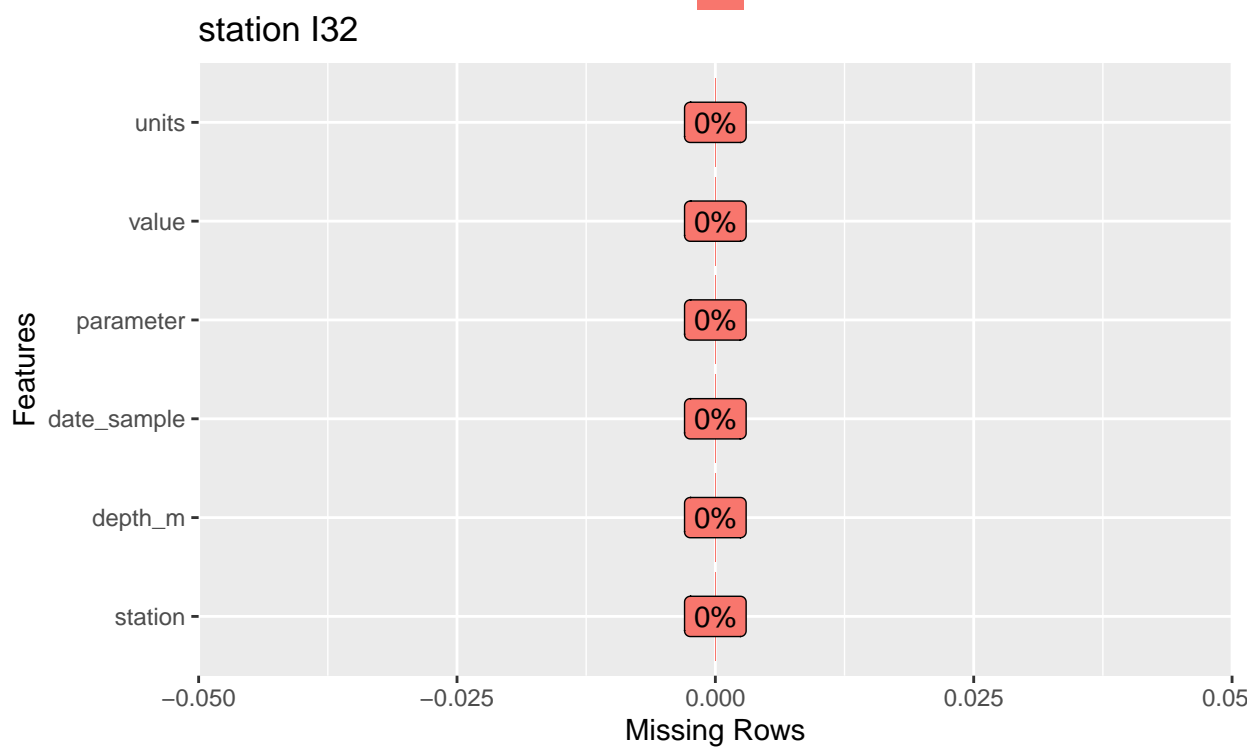
```









Band  Good



Band  Good



Band  Good



Band  Good

Date Sampling Deltas/Intervals for Dissolved Oxygen

Here we display the inconsistent sampling intervals which has led us into aggregating all stations as one, allowing for less likelihood of not having 4 samples per month.

```
# create I19 & I25 independent DFs for delta manipulation
I19 <- df %>%
  arrange(date_sample) %>%
  filter(station == "I19") %>%
  filter(parameter == "DO") %>%
  filter(date_sample >= "2011-01-01" & date_sample <= "2011-12-31") %>%
  select(-units)

I25 <- df %>%
  arrange(date_sample) %>%
  filter(station == "I25") %>%
  filter(parameter == "DO") %>%
  filter(date_sample >= "2011-01-01" & date_sample <= "2011-12-31") %>%
  select(-units)

I25_dates <- unique(I25$date_sample)
cat('\n', '-----I25 Date Deltas -----', '\n')

##
## -----I25 Date Deltas -----
I25_sampling_deltas <- (I25_dates[2:length(I25_dates)] - I25_dates)
I25_sampling_deltas[1:(length(I25_sampling_deltas)-1)]

## Time differences in days
## [1] 8 6 5 8 1 4 12 6 5 2 10 4 6 5 5 5 7 5 8 11 3 9 6 6 3
## [26] 5 6 6 6 10 7 6 6 6 6 4 9 4 7 7 8 4 5 5 3 4 7 6 6 11
## [51] 4 7 4 9 4 5 8 3 3

cat('\n\n')

I19_dates <- unique(I19$date_sample)
cat('\n', '-----I19 Date Deltas -----', '\n')

##
## -----I19 Date Deltas -----
I19_sampling_deltas <- (I19_dates[2:length(I19_dates)] - I19_dates)
I19_sampling_deltas[1:(length(I19_sampling_deltas)-1)]

## Time differences in days
## [1] 28 29 35 34 29 28 48 22 21 34 29
```

As we can see, DO for station I25 was sampled almost irregularly. With sampling intervals ranging from the sampling the next day to almost 2 weeks out versus the sampling frequency for I19 of once per month at irregular intervals of every 4-7 weeks.

Create time series data for each parameter, resampled to weekly observations, aggregated across all stations. 1 series per parameter

```
# Function to create time series from dataframe
create_ts <- function(parameter){
  p_ts <- wq_df[wq_df$parameter == parameter, c("datetime", "value")]
  p_ts_clean <- p_ts[!is.na(p_ts$datetime), ]
  # Order by datetime then convert to date
  p_ts_clean <- p_ts_clean[order(p_ts_clean$datetime), ]
  p_ts_clean$datetime <- as.Date(p_ts_clean$datetime)

  # Resample to weekly values by aggregating by week and taking the mean
  p_ts_rsmpl <- p_ts_clean %>%
    mutate(week = cut.Date(datetime, breaks = "1 week", labels = FALSE)) %>%
    group_by(week) %>%
    summarize(mean = mean(value, na.rm = TRUE))


  # set nan values to mean after resample
  p_ts_rsmpl$mean[is.na(p_ts_rsmpl$mean)] = mean(p_ts_clean$value, na.rm = TRUE)

  ret_ts <- ts(p_ts_rsmpl$mean, start=c(2011, 1), frequency=52)
  tsplot(ret_ts, main = paste("Weekly Mean ", parameter))

  return (ret_ts)
}
```

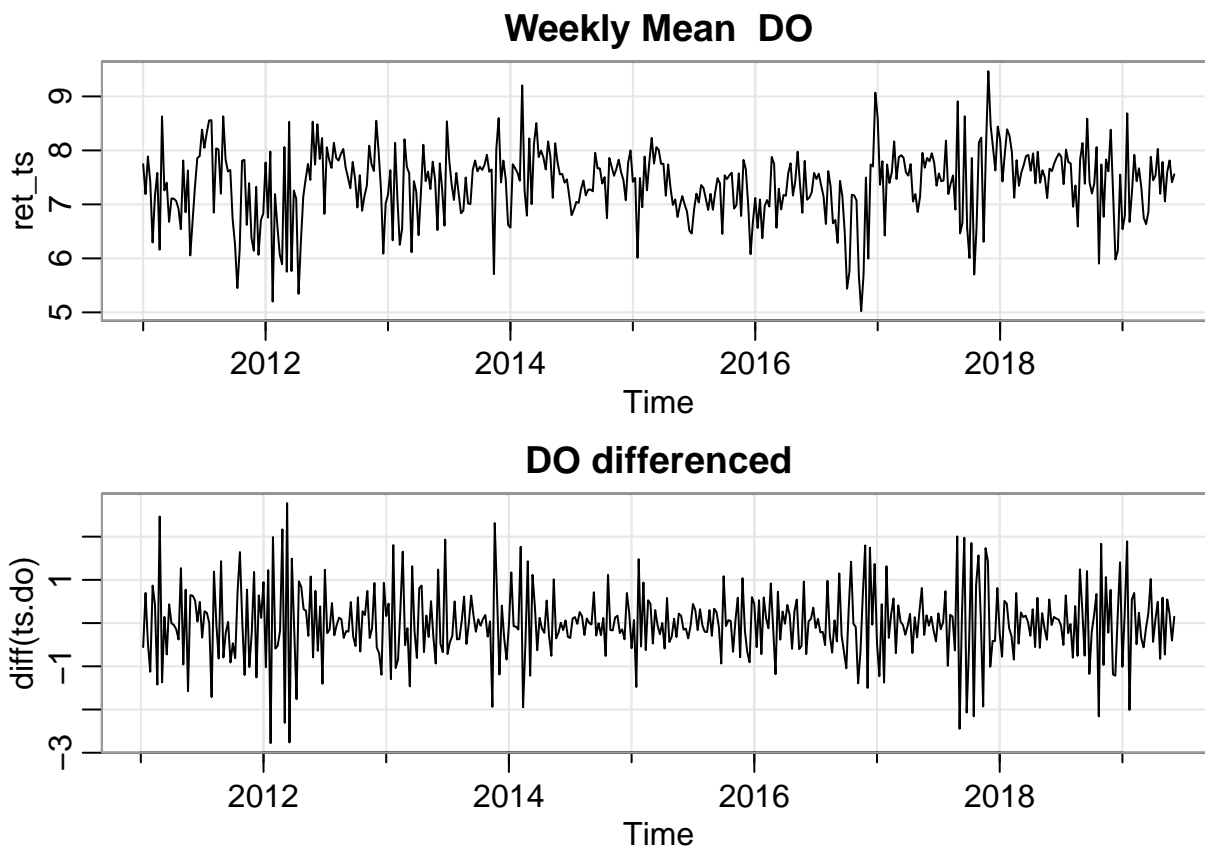
CHLOROPHYLL

```
par(mfrow=c(2, 1))
ts.chlorophyll <- create_ts('CHLOROPHYLL')
tsplot(diff(ts.chlorophyll), main = "CHLOROPHYLL differenced")
```




Dissolved Oxygen

```
par(mfrow=c(2, 1))  
ts.do <- create_ts('DO')  
tsplot(diff(ts.do), main = "DO differenced")
```




Entero

```
par(mfrow=c(2, 1))  
ts.entero <- create_ts('ENTERO')  
tsplot(diff(ts.entero), main = "ENTERO differenced")
```



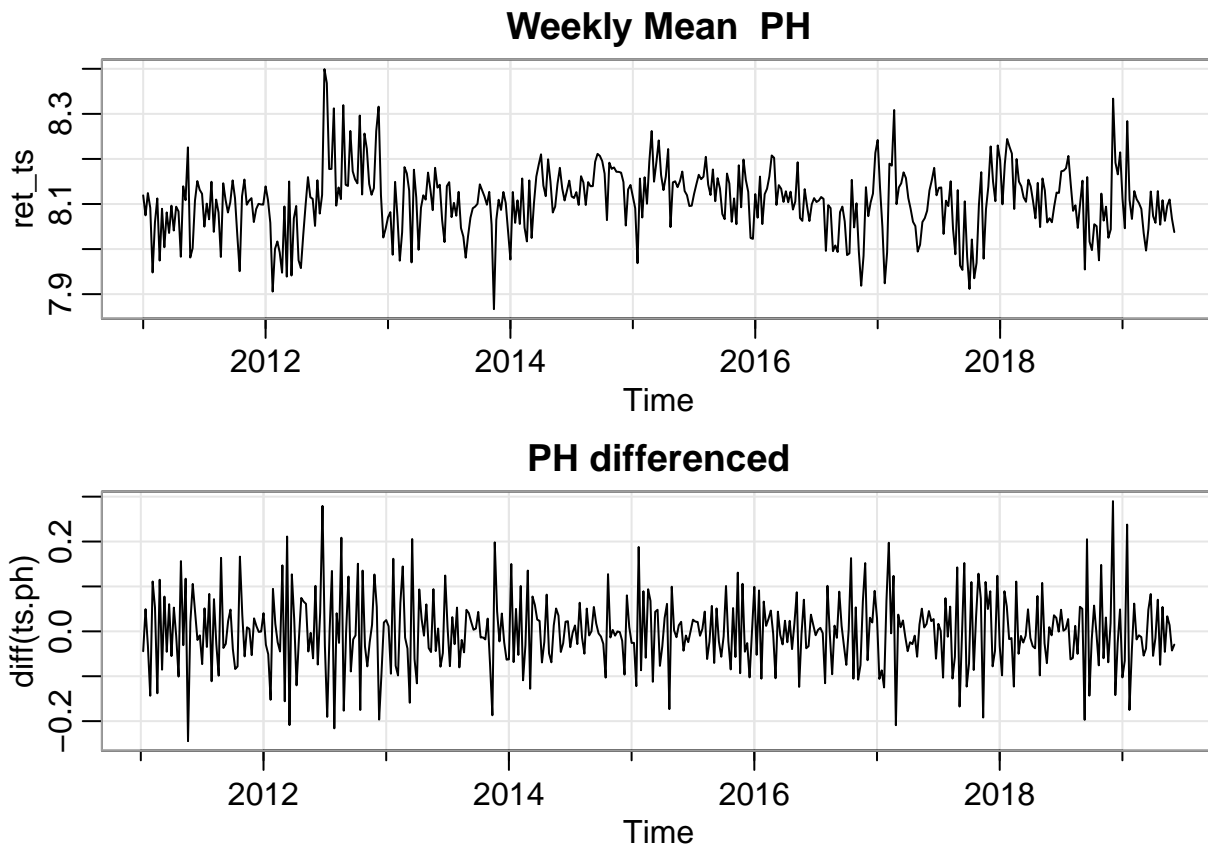
Fecal

```
par(mfrow=c(2, 1))
ts.fecal <- create_ts('FECAL')
tsplot(diff(ts.fecal), main = "FECAL differenced")
```


PH

```
par(mfrow=c(2, 1))
ts.ph <- create_ts('PH')
tsplot(diff(ts.ph), main = "PH differenced")
```



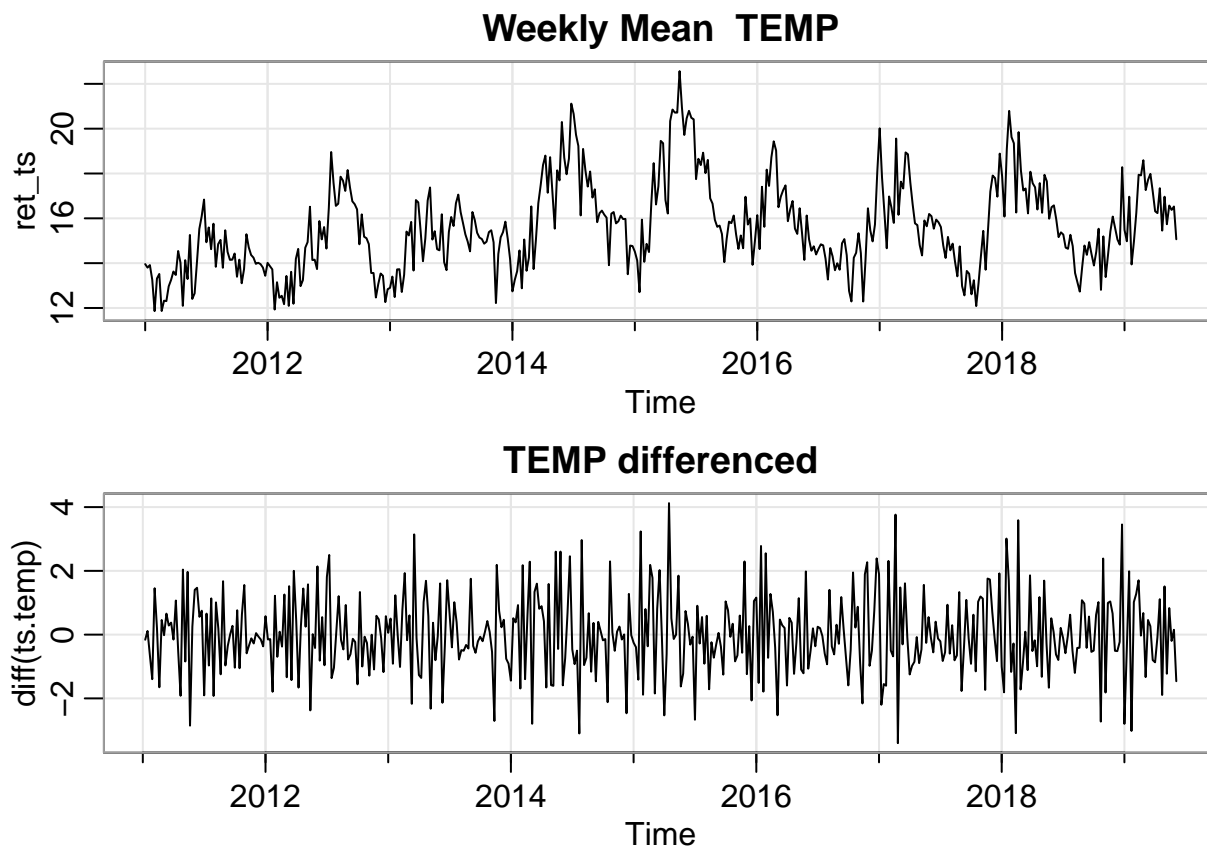
SALINITY

```
par(mfrow=c(2, 1))  
ts.salinity <- create_ts('SALINITY')  
tsplot(diff(ts.salinity), main = "SALINITY differenced")
```



TEMP

```
par(mfrow=c(2, 1))  
ts.temp <- create_ts('TEMP')  
tsplot(diff(ts.temp), main = "TEMP differenced")
```



ARIMA + ANN Model

```
# GENERALIZED Function to train ARIMA + ANN model
ARIMANN <- function(ts, forecast){
  # Keep things consistent
  set.seed(42)
  ts.size <- length(diff(ts))

  arima.model <- arima(diff(ts), order=c(1,1,1))
  arima.res <- arima.model$residuals

  # NN Window is 52 + 1 for label
  nn_train_set <- data.frame(matrix(ncol = 53, nrow = 0))

  for (i in 1:(ts.size - 53)) {
    nn_train_set <- rbind(nn_train_set, arima.res[i:(i+52)])
  }

  # Change label col name
  colnames(nn_train_set)[53] <- "Y"

  n <- names(nn_train_set)
  # For some reason R's neuralnet library can't properly parse a formula, we have to explicitly create
  f <- as.formula(paste("Y ~", paste(n[!n %in% "Y"], collapse = " + ")))
}
```

```

nn.model <- neuralnet(f, data = nn_train_set, linear.output = TRUE, learningrate = 0.01, hidden = 5)

# Check if forecast is blank
if (missing(forecast)){
  # Return training ts
  return (ts((diff(ts)[54:ts.size] - arima.res[54:ts.size]) + predict(nn.model, newdata = nn_train_set)
})

# Get the last 52 residuals to start the rolling predictions
nn.pred <- tail(arima.res, 52)
# iterate to forecast horizon for NN prediction
for (h in 1:forecast){
  pd.input <- data.frame(matrix(ncol = 52, nrow = 0))
  pd.input <- rbind(pd.input, tail(nn.pred, 52))
  colnames(pd.input) <- colnames(nn_train_set)[1:52]

  # append next prediction
  nn.pred <- append(nn.pred, predict(nn.model, newdata = pd.input))
}

pred.ts <- predict(arima.model, n.ahead=forecast)$pred + tail(nn.pred, forecast)

return (pred.ts)
}

```

Modeling Chlorophyll


```

# Only include data to 2018, reserve 2019 for validation
chlor.train <- window(ts.chlorophyll, 2011, c(2018, 52))
chlor.results <- ARIMANN(chlor.train, 23)
chlor.combined <- ts(c(diff(chlor.train), chlor.results), start=start(chlor.train), frequency = frequency(chlor.train))

# Invert the differencing
ts.plot(ts.chlorophyll, diffinv(chlor.combined, xi = ts.chlorophyll[1]), gpars = list(col = c("black", "red", "blue"),
abline(v=as.Date("2019-01-01"))

```

Chlorophy Prediction




RMSE

```
RMSE(tail(diffinv(chlor.combined, xi = ts.chlorophyll[1]),23), diff(window(ts.chlorophyll, 2019)))  
## [1] 6.351481
```

Modeling Dissolved Oxygen

```
# Only include data to 2018, reserve 2019 for validation  
do.train <- window(ts.do, 2011, c(2018, 52))  
do.results <- ARIMANN(do.train, 23)  
do.combined <- ts(c(diff(do.train), do.results), start=start(do.train), frequency = frequency(do.train))  
ts.plot(ts.do, diffinv(do.combined, xi = ts.do[1]), gpars = list(col = c("black", "red")), main = "Dissolved Oxygen")  
abline(v=as.Date("2019-01-01"), col = "blue")
```

Disolved Oxygen Prediction



RMSE

```
RMSE(tail(diffinv(do.combined, xi = ts.do[1]),23), diff(window(ts.do, 2019)))
```


```
## [1] 6.650365
```

Modeling Entero

```
# Only include data to 2018, reserve 2019 for validation
entero.train <- window(ts.entero, 2011, c(2018, 52))
entero.results <- ARIMANN(entero.train, 23)
entero.combined <- ts(c(diff(entero.train), entero.results), start=start(entero.train), frequency = frequency(entero.train))

ts.plot(window(ts.entero, 2011, c(2019, 23)), diffinv(entero.combined, xi = ts.entero[1]), gpars = list(lty=1, col="red"),
abline(v=as.Date("2019-01-01"), col = "blue")
```

Entero Prediction




RMSE

```
RMSE(tail(diffinv(entero.combined, xi = ts.entero[1]),23), diff(window(ts.entero, 2019)))  
## [1] 731.7727
```

Modeling Fecal

```
# Only include data to 2018, reserve 2019 for validation  
fecal.train <- window(ts.fecal, 2011, c(2018, 52))  
fecal.results <- ARIMANN(fecal.train, 23)  
fecal.combined <- ts(c(diff(fecal.train), fecal.results), start=start(fecal.train), frequency = frequency(fecal.train))  
  
ts.plot(window(ts.fecal, 2011, c(2019, 23)), diffinv(fecal.combined, xi = ts.fecal[1]), gpars = list(col = "red", lty = 1),  
abline(v=as.Date("2019-01-01"), col = "blue"))
```


Fecal Prediction




RMSE

```
RMSE(tail(diffinv(fecal.combined, xi = ts.fecal[1]), 23), diff(window(ts.fecal, 2019)))  
## [1] 1216.726
```

Modeling PH

```
# Only include data to 2018, reserve 2019 for validation  
ph.train <- window(ts.ph, 2011, c(2018, 52))  
ph.results <- ARIMANN(ph.train, 23)  
ph.combined <- ts(c(diff(ph.train), ph.results), start=start(ph.train), frequency = frequency(ph.train))  
  
ts.plot(window(ts.ph, 2011, c(2019, 23)), diffinv(ph.combined, xi = ts.ph[1]), gpars = list(col = c("blue", "red"),  
abline(v=as.Date("2019-01-01"), col = "blue")
```

PH Prediction



RMSE

```
RMSE(tail(diffinv(ph.combined, xi = ts.ph[1]), 23), diff(window(ts.ph, 2019)))
```


```
## [1] 8.17014
```

Modeling Salinity

```
# Only include data to 2018, reserve 2019 for validation
salinity.train <- window(ts.salinity, 2011, c(2018, 52))
salinity.results <- ARIMANN(salinity.train, 23)
salinity.combined <- ts(c(diff(salinity.train), salinity.results), start=start(salinity.train), frequency=frequency(salinity.train))

ts.plot(window(ts.salinity, 2011, c(2019, 23)), diffinv(salinity.combined, xi = ts.salinity[1]), gpars = list(lty = 1, col = "red"),
abline(v=as.Date("2019-01-01"), col = "blue"))
```

Salinity Prediction




RMSE

```
RMSE(tail(diffinv(salinity.combined, xi = ts.salinity[1]), 23), diff(window(ts.salinity, 2019)))  
## [1] 33.57443
```

Modeling Temp

```
# Only include data to 2018, reserve 2019 for validation  
temp.train <- window(ts.temp, 2011, c(2018, 52))  
temp.results <- ARIMANN(temp.train, 23)  
temp.combined <- ts(c(diff(temp.train), temp.results), start=start(temp.train), frequency = frequency(temp.train))  
  
ts.plot(window(ts.temp, 2011, c(2019, 23)), diffinv(temp.combined, xi = ts.temp[1]), gpars = list(col = "red", lty = 1),  
abline(v=as.Date("2019-01-01"), col = "blue"))
```

Temp Prediction



RMSE

```
RMSE(tail(diffinv(temp.combined, xi = ts.temp[1]), 23), diff(window(ts.temp, 2019)))

## [1] 19.5054
```

ARIMA Models


Chlorophyll

```
# Only include data to 2018, reserve 2019 for validation
chlor.arima <- auto.arima(window(ts.chlorophyll, 2011, c(2018, 52)))
summary(chlor.arima)


## Series: window(ts.chlorophyll, 2011, c(2018, 52))
## ARIMA(1,1,4)(1,0,0)[52]
##
## Coefficients:
##          ar1      ma1      ma2      ma3      ma4      sar1
##      -0.9455  0.2784 -0.8364 -0.2268  0.0144  0.0699
## s.e.   0.0485  0.0686  0.0580  0.0515  0.0495  0.0633
##
## sigma^2 estimated as 3.779:  log likelihood=-862.55
## AIC=1739.11  AICc=1739.38  BIC=1767.31
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
```

```
## Training set -0.05421516 1.927541 1.181362 -31.3413 52.54665 0.7195104
## ACF1
## Training set -0.002045334
```

```
par(mfrow = c(2,1))
plot(window(ts.chlorophyll, 2011, c(2019, 23)))
plot(forecast(chlor.arima, h = 23))
```



Forecasts from ARIMA(1,1,4)(1,0,0)[52]



RMSE

```
RMSE(forecast(chlor.arima, h = 23)$mean, window(ts.chlorophyll, 2019))
```

```
## [1] 1.042137
```


Dissolved Oxygen

```
# Only include data to 2018, reserve 2019 for validation
do.arima <- auto.arima(window(ts.do, 2011, c(2018, 52)))
summary(do.arima)
```


```
## Series: window(ts.do, 2011, c(2018, 52))
## ARIMA(2,0,2)(2,0,0)[52] with non-zero mean
##
## Coefficients:
##      ar1      ar2      ma1      ma2      sar1      sar2      mean
##      -0.0749  0.7711  0.2401 -0.5834  0.0282 -0.0496  7.3649
## s.e.    0.1034  0.0927  0.1216  0.1052  0.0531  0.0613  0.0656
##
## sigma^2 estimated as 0.4059: log likelihood=-399.47
## AIC=814.95 AICc=815.3 BIC=847.2
```

```
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.0007095884 0.6317496 0.4761292 -0.8049509 6.705114 0.689416
##           ACF1
## Training set 0.001939922

par(mfrow = c(2,1))
plot(window(ts.do, 2011, c(2019, 23)))
plot(forecast(do.arima, h = 23))
```



Forecasts from ARIMA(2,0,2)(2,0,0)[52] with non-zero mean



RMSE

```
RMSE(forecast(do.arima, h = 23)$mean, window(ts.do, 2019))
```

```
## [1] 0.5184086
```


Entero

```
# Only include data to 2018, reserve 2019 for validation
entero.arima <- auto.arima(window(ts.entero, 2011, c(2018, 52)))
summary(entero.arima)
```


```
## Series: window(ts.entero, 2011, c(2018, 52))
## ARIMA(3,0,1)(0,0,1)[52] with non-zero mean
##
## Coefficients:
##      ar1      ar2      ar3      ma1      sma1      mean
##      0.8963  0.0443 -0.0662 -0.7547  0.1155  89.1923
## s.e.  0.1776  0.0819  0.0755  0.1656  0.0612  25.0402
##
## sigma^2 estimated as 56298: log likelihood=-2862.87
```

```
## AIC=5739.73   AICc=5740.01   BIC=5767.95
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -5.198707 235.5545 107.4697 -1016.96 1035.669 0.9052769
##           ACF1
## Training set -0.04759006

par(mfrow = c(2,1))
plot(window(ts.entero, 2011, c(2019, 23)))
plot(forecast(entero.arima, h = 23))
```



Forecasts from ARIMA(3,0,1)(0,0,1)[52] with non-zero mean



RMSE

```
RMSE(forecast(do.arima, h = 23)$mean, window(ts.entero, 2019))
```

```
## [1] 467.3004
```


Fecal

```
# Only include data to 2018, reserve 2019 for validation
fecal.arima <- auto.arima(window(ts.fecal, 2011, c(2018, 52)))
summary(fecal.arima)
```


```
## Series: window(ts.fecal, 2011, c(2018, 52))
## ARIMA(1,0,1)(1,0,0)[52] with non-zero mean
##
## Coefficients:
##           ar1      ma1      sar1      mean
##           0.5660 -0.3393  0.1838 106.5577
## s.e.      1.0108  1.1306  0.0905  37.7856
##
```

```
## sigma^2 estimated as 127140: log likelihood=-3033.84
## AIC=6077.67 AICc=6077.82 BIC=6097.82
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set -9.167291 354.8481 135.2333 -1358.5 1380.195 1.05456 -0.1425296

par(mfrow = c(2,1))
plot(window(ts.fecal, 2011, c(2019, 23)))
plot(forecast(fecal.arima, h = 23))
```



Forecasts from ARIMA(1,0,1)(1,0,0)[52] with non-zero mean



RMSE

```
RMSE(forecast(fecal.arima, h = 23)$mean, window(ts.fecal, 2019))
```

```
## [1] 567.3506
```

PH


```
# Only include data to 2018, reserve 2019 for validation
ph.arima <- auto.arima(window(ts.ph, 2011, c(2018, 52)))
summary(ph.arima)
```

```
## Series: window(ts.ph, 2011, c(2018, 52))
## ARIMA(1,0,2)(0,0,1)[52] with non-zero mean
##
## Coefficients:
##          ar1          ma1          ma2          sma1          mean
##          0.8849    -0.5778    -0.0905     0.0907     8.1083
## s.e.        0.0462     0.0686     0.0560     0.0507     0.0100
```




```
##
## sigma^2 estimated as 0.004497: log likelihood=535.97
## AIC=-1059.94 AICc=-1059.73 BIC=-1035.75
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.0001861881 0.0666553 0.05046257 -0.004469943 0.6227442 0.6627212
##           ACF1
## Training set -0.00308448

par(mfrow = c(2,1))
plot(window(ts.ph, 2011, c(2019, 23)))
plot(forecast(ph.arima, h = 23))
```



Forecasts from ARIMA(1,0,2)(0,0,1)[52] with non-zero mean



RMSE

```
RMSE(forecast(ph.arima, h = 23)$mean, window(ts.ph, 2019))
```

```
## [1] 0.06198158
```


Salinity

```
# Only include data to 2018, reserve 2019 for validation
salinity.arima <- auto.arima(window(ts.salinity, 2011, c(2018, 52)))
summary(salinity.arima)
```


```
## Series: window(ts.salinity, 2011, c(2018, 52))
## ARIMA(2,1,4)(2,0,0)[52] with drift
##
## Coefficients:
```

```
##          ar1      ar2      ma1      ma2      ma3      ma4      sar1      sar2      drift
##      -1.6533  -0.7358  1.3445  0.0797  -0.3442  0.0010  0.0424  0.0523  0.0011
## s.e.   0.2006   0.1591  0.2052  0.1383   0.1195  0.0612  0.0560  0.0592  0.0018
##
## sigma^2 estimated as 0.003259:  log likelihood=603.73
## AIC=-1187.46   AICc=-1186.92   BIC=-1147.18
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
## Training set 6.855504e-05 0.05639824 0.04131292 1.535374e-05 0.1233572
##              MASE      ACF1
## Training set 0.3319615 -0.0004661822
```

```
par(mfrow = c(2,1))
plot(window(ts.salinity, 2011, c(2019, 23)))
plot(forecast(salinity.arima, h = 23))
```



Forecasts from ARIMA(2,1,4)(2,0,0)[52] with drift



RMSE

```
RMSE(forecast(ph.arima, h = 23)$mean, window(ts.ph, 2019))
```

```
## [1] 0.06198158
```


Temp

```
# Only include data to 2018, reserve 2019 for validation
temp.arima <- auto.arima(window(ts.temp, 2011, c(2018, 52)))
summary(temp.arima)
```


```
## Series: window(ts.temp, 2011, c(2018, 52))
```

```
## ARIMA(0,1,2)(0,0,1)[52]
##
## Coefficients:
##          ma1      ma2      sma1
##       -0.4368 -0.0679  0.0863
## s.e.    0.0493   0.0454  0.0490
##
## sigma^2 estimated as 1.34:  log likelihood=-648.38
## AIC=1304.77   AICc=1304.86   BIC=1320.88
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.01166253 1.151942 0.8978906 -0.3194333 5.76851 0.5398124
##              ACF1
## Training set -0.001339407
```

```
par(mfrow = c(2,1))
plot(window(ts.temp, 2011, c(2019, 23)))
plot(forecast(temp.arima, h = 23))
```



Forecasts from ARIMA(0,1,2)(0,0,1)[52]



RMSE

```
RMSE(forecast(temp.arima, h = 23)$mean, window(ts.temp, 2019))
```

```
## [1] 1.205087
```